Westminster College: A Retention Prediction Model

Ryan P. Smith

Northwest Missouri State University, Maryville MO 64468, USA
S547012@nwmissouri.edu

Abstract. The abstract should briefly summarize the contents of the paper in

Keywords: higher education · retention · data analytics · data science

1 Introduction

At Westminster College, the Office of Institutional Effectiveness is responsible for compiling and analyzing data at our institution. This project will focus on the domain of higher education, and more specifically, we will be examining retention. The domain of higher education is being examined in this project as it can be utilized within our department in the future, and hopefully help provide more insight into retention factors at our college.

The data for this project will be sourced from Westminster College. We will be utilizing a couple of our databases from both our Enrollment Customer Relationship Management tool (CRM) Slate by Tehcnolutions, and our main SIS (Student Information System) which is provided by Jenzabar. All the student data used will be anonymized to ensure that there is no personal identifiable information used. The data will be compiled from the many tables into a useable dataset for this project. Data features that will be included in the dataset will be some biographical such as gender, commuter/resident, athlete. Some other data features will be items such as ACT score and high school GPA. This project will look to take those factors, create a multi-variable regression model to predict the likelihood that a student at Westminster College will be retained or not.

1.1 Defining the Problem

For this project, we want to see if retention can be predicted at Westminster College based on the selected data points and a machine learning model. There are multiple reasons why this is an important issue to look at and examine. First there appears to be a shrinking pool of high school graduates as we approach 2030. [5] This means that if there are fewer students in the graduating high school classes to recruit, there will be more competition between colleges and universities to enroll these students. If this leads to lower enrollment at an institution, that means that retention will become even more important to retain the students that you are able to enroll.

We can also see that enrollment has currently been dropping overall for the past several years. [4] The National Student Clearinghouse also illustrates this. We can see the enrollment declines among the different higher educations types, with a slight uptick actually showing in community colleges. [2] Since Westminster College has two prominent sources of revenue, donations and enrollment, it will be important to retain students at a higher percentage than before, if the enrollment of students continues to decline.

1.2 Goals of this Research and Methods to explore

This project will be looking to successfully predict if a student is at risk of being retained. We can use this data with our Student Success Center so that they can set up an intervention, or take the appropriate actions to ensure they can help the student persist. After gathering all the data to analyze, a model will be built and trained on train data, and then tested. After the model is completed, we can discuss the results and see how the model performed.

This project will be following the traditional Data Science Life Cycle. The steps are outlined below as:

1. Business Understanding - as seen above, this will be Predicting Retention at Westminster College.
2. Data Collecting - Data set to be produced from our CRM and SIS.
3. Data Cleaning - Missing data will be handled as well as mismatched data types.
4. Exploratory Data Analysis - Here we will begin to find trends and groups in the data.
5. Model Building - Here the model method will be selected, and then trained and tested.
6. Results - Visualizations will be created, conclusions will be formed and discussed.

The key components that will be focused on for this project will be data from previously enrolled students. We will be focusing only on students that attended Westminster College, and utilize the data that we have access to for each student. Some items may be missing such as ACT from recent years. Enrollment recently went to ACT optional, so there are more missing ACT scores than from the past, as more students are choosing not to submit those for consideration. We will look at including and not including ACT scores as another trend has been the decline in overall ACT scores across the country [3], and the state of Missouri [1] (19.5 and 19.8 respectively). This could be an interesting limitation, so we will try looking at retention with and without this feature.

2 About the Data

Data was collected, cleaned, and analyzed for this project. It came directly from two databases from Westminster College. These two databases come from the

Student Information System (SIS) which is provided by Jenzabar. The other is our Enrollment Customer Relationship Management (CRM) called Slate, by Technolutions. The data was pulled directly from a combination of these two databases, compiled into a usable dataset within a csv file, cleaned using python, all before being analyzed.

2.1 Data Sourcing

The data was sourced from two databases on campus, the SIS and our Enrollment CRM. This data is structured data stored within these two databases. All tables being used to source our data are able to be joined together based on the student ID. The data was sourced directly from the databases, and the results were written to a .csv file to be imported into a python script for cleaning and initial data exploration. There was no scraping needed to gather this data, it was collected with a SQL script using SQL Server. All tables and the two databases were able to be joined together with the primary key of student ID.

When looking back at students entering from the year 2010, there are 2605 records in the data set, when expanded to the year 2000, the amount of records increased to 5265. This project utilized the latter to have more records for analysis. For this project, the data fields were limited to 7 total. The fields in the data set along with their respective data type will be as follows:

1. Unique Identifier: Numeric Integer - Used to identify a unique student within the data set.
2. ACT Score: Numeric Integer - The highest ACT score that a student submitted with their enrollment.
3. Resident/Commuter: Character - denotes is a student resides on campus or commutes to campus
4. Gender: Character
5. Athlete: String - A sport will denote whether the student is an athlete or not (converted to a Y or N)
6. High School GPA: Numeric Float - a float numeric value to show the GPA from the students graduating High School.
7. Graduation Data/Exit Date: date - Currently a data field for both columns, these were combined along with exit reason to establish if a student graduated and was retained or withdrew and was not retained.

Overall, the data is very clean. There are some missing values in ACT scores, which was expected since we became an ACT optional reporting school in terms of enrollment. Looking through the data there are several other missing values, but not enough to have to cut it out. During the cleaning phase those issues will be addressed. There are two date fields that are formatted as YYYY,mm,dd, however, those fields were combined into a single field to determine if the student was retained (graduated) or not (withdrew). There were no bad characters found in the data set.

All of the data compiled from the databases was executed and written to a .csv document. With this format, the data was taken and using various python libraries, cleaned and explored (both examined more in the upcoming sections). Also, as mentioned there are a couple items where were combined from the elements from the database to make a single feature within the dataset. After collecting and compiling all the data, the next section will cover how the data was cleaned.

2.2 Cleaning the Data

After collecting all the data, it was exported from the SQL Server environment into a .csv file. A python script had been written to pull the .csv file into a pandas data frame for further analysis. After the data was imported into a dataframe, the dataframe was printed and also looked at the columns, and described the data. That can be seen in the image figure 1.

When looking at the data itself the cleaning looked minimal. However, using Python there were a few things discovered that needed to be cleaned. The first cleaning that was discovered was the duplicate values. The though going in was there would be no duplicates, however, there were some as the ACT test switched formats and some students had taken both tests, and had two highest scores recoreded. The script that was used to remove those were:

$$data.drop_duplicates(inplace = True) \tag{1}$$

Once the duplicates were dropped, there were 5063 records remaining in the data set still with 7 the features.After getting the duplicates removed from the data set, the next thing to check was the null values. There were two columns that contained null values, and they were ACT Scores and High School GPA. A couple of scripts were ran against the datafame to see what columns contained the null values, as well as what percent of our data had a null in each of those columns. They can be seen in figure 2.

Below is a table showing the results of the null values in the data. Again we can see the only two columns were ACT Score and High School GPA. They had 17.28 percent and 22.71 percent respectively in null values.

Table 1. Null Data Percentage by Feature

Feature	Percent Null
ID NUM	0.00%
Gender	0.00%
Res Communter Sts	0.00%
ACT SCORE	17.27%
ACT Score Received	0.00%
HS GPA	22.71%
Athlete	0.00%
Retained	0.00%

After identifying the null data, it was decided to use the mean value to fill in the missing values for both the ACT Scores and the High School GPA. The mean was selected as there was not a wide range of outliers being that ACT Scores and GPAs are relatively narrow and restricted in range. The script to replace the null values can be seen below.

```
data.act_scores.fillna(data.act_scores.mean()).isnull().any()
```

Once the two columns had their null values replaced with the mean, all the columns data was complete. The duplicates were removed, the nulls were replaced with the mean of their respective columns, and there were no bad characters in any column. Below we can see the features that will be analyzed after the cleaning of the data.

Table 2. Features Used After Data Cleaning

Feature	Type
ID NUM	Int
Gender	Character
Res Communter Sts	Character
ACT SCORE	float
ACT Score Received	Character
HS GPA	float
Athlete	character
Retained	character

The features themselves are defined as below.

1. Unique Identifier: Used to identify a unique student within the data set.
2. ACT Score: The highest ACT score that a student submitted with their enrollment.
3. Resident/Commuter: denotes is a student resides on campus or commutes to campus
4. Gender: A Students self reported gender
5. Athlete: Whether or not a student is an athlete
6. High School GPA: gpa from high school
7. Retained: Whether or not a student graduated and was retained at the college.

After cleaning the data we will be using 6 of the features above (ID number is just an identifier) along with the 5063 entries to analyze the data. The data will be used to see if it can predict whether a student will be retained or not at Westminster College. The dependent variable is going to be whether or not a student was retained through gradation at the college. We will use the other features of ACT Score/ACT Provided, Gender, Athlete status, High school GPA, and Resident Status to make that prediction and they will be the independent variables. ACT Scores and ACT Provided will be used in separate models to see if scores affect retention, or if having a score at all is effective. The next section will show some basic discover and exploration of the data.

3 Model Building

3.1 Choosing the Model

Decide which model to use with our data.

3.2 Building the Model

Walk through the building of the chosen model.

3.3 Training and Testing the Model

Create a training and test data set to use with the model.

4 Results

4.1 Clear Summary of Results

What did we find out from our model?

4.2 Visualizations

Share some visuals from the results and show what the data says.

5 Discussion

5.1 Conclusion of Results

5.2 Limitations

5.3 Future Recommendations and Work

References

1. Act scores drop for 6th straight year, https://www.northwestmoinfo.com/local-news/act-scores-drop-for-6th-straight-year/
2. Current term enrollment estimates, https://nscresearchcenter.org/current-term-enrollment-estimates/
3. Castillo, E.: Act scores hit 30-year low, https://www.bestcolleges.com/news/act-test-scores-hit-30-year-low
4. Knox, L.: Leveling off on the bottom, https://www.insidehighered.com/news/admissions/traditional-age/2023/05/24/leveling-bottom
5. Seltzer, R.: Birth dearth approaches, https://www.insidehighered.com/news/2020/12/15/more-high-school-graduates-through-2025-pool-still-shrinks-afterward

```
[53]: print(data)

      ID_NUM GENDER RES_COMMUTER_STS  ACT_SCORE ACT_RECEIVED  HS_GPA Athlete  \
0          1      F                R       22.0            Y    3.21       N
1          2      F                R       22.0            Y    3.66       Y
2          3      M                R        NaN            N    2.88       N
3          4      F                R       30.0            Y    3.73       Y
4          5      F                R        NaN            N    3.92       Y
...      ...    ...              ...        ...          ...     ...     ...
5059    5060      F                C        NaN            N     NaN       N
5060    5061      F                C       17.0            Y     NaN       N
5061    5062      F                R        NaN            N     NaN       N
5062    5063      M                R        NaN            N     NaN       Y
5063    5064      M                R       18.0            Y     NaN       Y

     RETAINED
0           N
1           Y
2           Y
3           Y
4           Y
...       ...
5059        Y
5060        Y
5061        N
5062        Y
5063        Y

[5064 rows x 8 columns]
```

```
[54]: data.columns

[54]: Index(['ID_NUM', 'GENDER', 'RES_COMMUTER_STS', 'ACT_SCORE', 'ACT_RECEIVED',
             'HS_GPA', 'Athlete', 'RETAINED'],
            dtype='object')
```

```
[55]: data.describe()
```

[55]:

	ID_NUM	ACT_SCORE	HS_GPA
count	5064.000000	4189.000000	3914.000000
mean	2532.500000	23.889472	2.963781
std	1461.995212	4.022140	1.298116
min	1.000000	11.000000	0.000000
25%	1266.750000	21.000000	2.830000
50%	2532.500000	24.000000	3.400000
75%	3798.250000	27.000000	3.800000
max	5064.000000	35.000000	5.140000

Fig. 1. Looking at data in the dataframe, displaying the columns, and a basic describe

```
[45]: data.isnull().any()

[45]: ID_NUM               False
      GENDER               False
      RES_COMMUTER_STS     False
      ACT_SCORE             True
      ACT_RECEIVED         False
      HS_GPA                True
      Athlete              False
      RETAINED             False
      dtype: bool

[46]: data.isnull().sum()/data.shape[0]

[46]: ID_NUM               0.000000
      GENDER               0.000000
      RES_COMMUTER_STS     0.000000
      ACT_SCORE            0.172788
      ACT_RECEIVED         0.000000
      HS_GPA               0.227093
      Athlete              0.000000
      RETAINED             0.000000
      dtype: float64

[64]: #get rid of the null values

      data.ACT_SCORE.fillna(data.ACT_SCORE.mean())

[64]: 0       22.000000
      1       22.000000
      2       23.889472
      3       30.000000
      4       23.889472
                ...
      5059    23.889472
      5060    17.000000
      5061    23.889472
      5062    23.889472
      5063    18.000000
      Name: ACT_SCORE, Length: 5064, dtype: float64
```

Fig. 2. Columns that show null data, and the percent of each that are null within each one.